

With TF 1.0!



Lab 10

NN, ReLu, Xavier, Dropout, and Adam

Sung Kim <hunkim+ml@gmail.com>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



Call for comments

Please feel free to add comments directly on these slides

Other slides: <https://goo.gl/jPtVNT>



With TF 1.0!



Lab 10

NN, ReLu, Xavier, Dropout, and Adam

Sung Kim <hunkim+ml@gmail.com>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/>



zeran4

1 commit / 5 ++ / 4 --

#11



jennykang

19 commits / 940 ++ / 253 --

#2



GzuPark

14 commits / 41 ++ / 31 --

#3



kkweon

5 commits / 372 ++ / 296 --

#4



BlueMelon715

4 commits / 45 ++ / 34 --

#5



jin-chong

2 commits / 4 ++ / 4 --

#6



FuZer

2 commits / 37 ++ / 30 --

#7



cynthia

1 commit / 28 ++ / 28 --

#8



keon

1 commit / 3 ++ / 3 --

#9



allieus

1 commit / 55 ++ / 59 --

#10

Softmax classifier for MNIST

```
# weights & bias for nn layers inits outputs
W = tf.Variable(tf.random_normal([784, 10]))
b = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(X, W) + b
# define cost/loss & optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=hypothesis, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
# initialize
sess = tf.Session()
sess.run(tf.global_variables_initializer())

# train my model
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = int(mnist.train.num_examples / batch_size)

    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
        feed_dict = {X: batch_xs, Y: batch_ys}
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)
        avg_cost += c / total_batch

    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))

print('Learning Finished!')

# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy, feed_dict={X: mnist.test.images, Y: mnist.test.labels}))
```

```
Epoch: 0001 cost = 5.888845987
Epoch: 0002 cost = 1.860620173
Epoch: 0003 cost = 1.159035648
Epoch: 0004 cost = 0.892340870
Epoch: 0005 cost = 0.751155428
Epoch: 0006 cost = 0.662484806
Epoch: 0007 cost = 0.601544010
Epoch: 0008 cost = 0.556526115
Epoch: 0009 cost = 0.521186961
Epoch: 0010 cost = 0.493068354
Epoch: 0011 cost = 0.469686249
Epoch: 0012 cost = 0.449967254
Epoch: 0013 cost = 0.433519321
Epoch: 0014 cost = 0.419000337
Epoch: 0015 cost = 0.406490815
Learning Finished!
Accuracy: 0.9035
```

With TF 1.0!

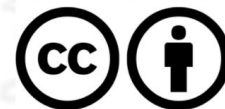


Lab 7-2

MNIST data

Sung Kim <hunkim+ml@gmail.com>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



Softmax classifier for MNIST

```
# weights & bias for nn layers
W = tf.Variable(tf.random_normal([784, 10]))
b = tf.Variable(tf.random_normal([10]))
hypothesis = tf.matmul(X, W) + b
# define cost/loss & optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=hypothesis, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
# initialize
sess = tf.Session()
sess.run(tf.global_variables_initializer())

# train my model
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = int(mnist.train.num_examples / batch_size)

    for i in range(total_batch):
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
        feed_dict = {X: batch_xs, Y: batch_ys}
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)
        avg_cost += c / total_batch

    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))

print('Learning Finished!')

# Test model and check accuracy
correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
print('Accuracy:', sess.run(accuracy, feed_dict={X: mnist.test.images, Y: mnist.test.labels}))
```

```
Epoch: 0001 cost = 5.888845987
Epoch: 0002 cost = 1.860620173
Epoch: 0003 cost = 1.159035648
Epoch: 0004 cost = 0.892340870
Epoch: 0005 cost = 0.751155428
Epoch: 0006 cost = 0.662484806
Epoch: 0007 cost = 0.601544010
Epoch: 0008 cost = 0.556526115
Epoch: 0009 cost = 0.521186961
Epoch: 0010 cost = 0.493068354
Epoch: 0011 cost = 0.469686249
Epoch: 0012 cost = 0.449967254
Epoch: 0013 cost = 0.433519321
Epoch: 0014 cost = 0.419000337
Epoch: 0015 cost = 0.406490815
Learning Finished!
Accuracy: 0.9035
```

NN for MNIST

```
# input place holders
```

```
X = tf.placeholder(tf.float32, [None, 784])
```

```
Y = tf.placeholder(tf.float32, [None, 10])
```

```
# weights & bias for nn layers
```

```
W1 = tf.Variable(tf.random_normal([784, 256]))
```

```
b1 = tf.Variable(tf.random_normal([256]))
```

```
L1 = tf.nn.relu(tf.matmul(X, W1) + b1)
```

```
W2 = tf.Variable(tf.random_normal([256, 256]))
```

```
b2 = tf.Variable(tf.random_normal([256]))
```

```
L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)
```

```
W3 = tf.Variable(tf.random_normal([256, 10]))
```

```
b3 = tf.Variable(tf.random_normal([10]))
```

```
hypothesis = tf.matmul(L2, W3) + b3
```

```
# define cost/loss & optimizer
```

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
```

```
    logits=hypothesis, labels=Y))
```

```
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

Epoch: 0001 cost = 141.207671860

Epoch: 0002 cost = 38.788445864

Epoch: 0003 cost = 23.977515479

Epoch: 0004 cost = 16.315132428

Epoch: 0005 cost = 11.702554882

Epoch: 0006 cost = 8.573139748

Epoch: 0007 cost = 6.370995680

Epoch: 0008 cost = 4.537178684

Epoch: 0009 cost = 3.216900532

Epoch: 0010 cost = 2.329708954

Epoch: 0011 cost = 1.715552875

Epoch: 0012 cost = 1.189857912

Epoch: 0013 cost = 0.820965160

Epoch: 0014 cost = 0.624131458

Epoch: 0015 cost = 0.454633765

Learning Finished!

Accuracy: 0.9455



All

Videos

News

Images

Maps

More

Settings

Tools

About 2,000 results (0.43 seconds)

python - How to do Xavier initialization on TensorFlow - Stack Overflow

stackoverflow.com/questions/33640581/how-to-do-xavier-initialization-on-tensorflow ▾

Nov 10, 2015 - I'm porting my Caffe network over to TensorFlow but it doesn't seem to ... Does Xavier initialization changes how the biases are initialized ...

You've visited this page many times. Last visit: 1/26/17

초기화를 잘하는법

n

Tags

Users

Search...

on TensorFlow

to TensorFlow but it doesn't seem to have xavier initialization. I'm seems to be making it a lot harder to train.

edited Dec 19 '15 at 22:25



Hooked

32.2k ● 13 ● 108 ● 169

asked Nov 10 '15 at 22:07



Aleph7

2,646 ● 1 ● 14 ● 24

Does Xavier initialization changes how the biases are initialized? – Pinocchio Jul 25 '16 at 20:07

[add a comment](#)[start a bounty](#)

4 Answers

active

oldest

votes



Since version 0.8 there is a Xavier initializer, [see here for the docs](#).

55

You can use something like this:



```
W = tf.get_variable("W", shape=[784, 256],  
                    initializer=tf.contrib.layers.xavier_initializer())
```

[share](#) [edit](#) [delete](#) [flag](#)

edited Dec 1 '16 at 16:32



fabian789

5,296 ● 3 ● 32 ● 75

answered Apr 22 '16 at 4:23



Sung Kim

2,339 ● 1 ● 15 ● 24

Xavier for MNIST

input place holders

X = tf.placeholder(tf.float32, [None, 784])

Y = tf.placeholder(tf.float32, [None, 10])

weights & bias for nn layers

<http://stackoverflow.com/questions/33640581>

W1 = tf.get_variable("W1", shape=[784, 256],
initializer=tf.contrib.layers.xavier_initializer())

b1 = tf.Variable(tf.random_normal([256]))

L1 = tf.nn.relu(tf.matmul(X, W1) + b1)

W2 = tf.get_variable("W2", shape=[256, 256],
initializer=tf.contrib.layers.xavier_initializer())

b2 = tf.Variable(tf.random_normal([256]))

L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)

W3 = tf.get_variable("W3", shape=[256, 10],
initializer=tf.contrib.layers.xavier_initializer())

b3 = tf.Variable(tf.random_normal([10]))

hypothesis = tf.matmul(L2, W3) + b3

define cost/loss & optimizer

cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
logits=hypothesis, labels=Y))

optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)

→ 초기화 방식은 바꿀

→ 처음부터 cost가 낮다.

Epoch: 0001 cost = 0.301498963
Epoch: 0002 cost = 0.107252513
Epoch: 0003 cost = 0.064888892
Epoch: 0004 cost = 0.044463030
Epoch: 0005 cost = 0.029951642
Epoch: 0006 cost = 0.020663404
Epoch: 0007 cost = 0.015853033
Epoch: 0008 cost = 0.011764387
Epoch: 0009 cost = 0.008598264
Epoch: 0010 cost = 0.007383116
Epoch: 0011 cost = 0.006839140
Epoch: 0012 cost = 0.004672963
Epoch: 0013 cost = 0.003979437
Epoch: 0014 cost = 0.002714260
Epoch: 0015 cost = 0.004707661
Learning Finished!

Accuracy: **0.9783**

Xavier for MNIST

초깃값이 잘 차라
되잖아.

```
# input place holders
X = tf.placeholder(tf.float32, [None, 784])
Y = tf.placeholder(tf.float32, [None, 10])
```

```
# weights & bias for nn layers
```

```
# http://stackoverflow.com/questions/3364
```

```
W1 = tf.get_variable("W1", shape=[784, 256],  
                    initializer=tf.contrib
```

```
b1 = tf.Variable(tf.random_normal([256]))
```

```
L1 = tf.nn.relu(tf.matmul(X, W1) + b1)
```

```
W2 = tf.get_variable("W2", shape=[256, 256],  
                    initializer=tf.contrib
```

```
b2 = tf.Variable(tf.random_normal([256]))
```

```
L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)
```

```
W3 = tf.get_variable("W3", shape=[256, 10],  
                    initializer=tf.contrib
```

```
b3 = tf.Variable(tf.random_normal([10]))
```

```
hypothesis = tf.matmul(L2, W3) + b3
```

```
# define cost/loss & optimizer
```

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(  
    logits=hypothesis, labels=Y))
```

```
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

```
Epoch: 0001 cost = 141.207671860
Epoch: 0002 cost = 38.788445864
Epoch: 0003 cost = 23.977515479
Epoch: 0004 cost = 16.315132428
Epoch: 0005 cost = 11.702554882
Epoch: 0006 cost = 8.573139748
Epoch: 0007 cost = 6.370995680
Epoch: 0008 cost = 4.537178684
Epoch: 0009 cost = 3.216900532
Epoch: 0010 cost = 2.329708954
Epoch: 0011 cost = 1.715552875
Epoch: 0012 cost = 1.189857912
Epoch: 0013 cost = 0.820965160
Epoch: 0014 cost = 0.624131458
Epoch: 0015 cost = 0.454633765
Learning Finished!
Accuracy: 0.9455 (normal dist)
```

```
Epoch: 0001 cost = 0.301498963
Epoch: 0002 cost = 0.107252513
Epoch: 0003 cost = 0.064888892
Epoch: 0004 cost = 0.044463030
Epoch: 0005 cost = 0.029951642
Epoch: 0006 cost = 0.020663404
Epoch: 0007 cost = 0.015853033
Epoch: 0008 cost = 0.011764387
Epoch: 0009 cost = 0.008598264
Epoch: 0010 cost = 0.007383116
Epoch: 0011 cost = 0.006839140
Epoch: 0012 cost = 0.004672963
Epoch: 0013 cost = 0.003979437
Epoch: 0014 cost = 0.002714260
Epoch: 0015 cost = 0.004707661
Learning Finished!
Accuracy: 0.9783 (xavier)
```

Deep NN for MNIST

and wide

```
W1 = tf.get_variable("W1", shape=[784, 512],  
    initializer=tf.contrib.layers.xavier_initializer())  
b1 = tf.Variable(tf.random_normal([512]))  
L1 = tf.nn.relu(tf.matmul(X, W1) + b1)  
  
W2 = tf.get_variable("W2", shape=[512, 512],  
    initializer=tf.contrib.layers.xavier_initializer())  
b2 = tf.Variable(tf.random_normal([512]))  
L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)  
  
W3 = tf.get_variable("W3", shape=[512, 512],  
    initializer=tf.contrib.layers.xavier_initializer())  
b3 = tf.Variable(tf.random_normal([512]))  
L3 = tf.nn.relu(tf.matmul(L2, W3) + b3)  
  
W4 = tf.get_variable("W4", shape=[512, 512],  
    initializer=tf.contrib.layers.xavier_initializer())  
b4 = tf.Variable(tf.random_normal([512]))  
L4 = tf.nn.relu(tf.matmul(L3, W4) + b4)  
  
W5 = tf.get_variable("W5", shape=[512, 10],  
    initializer=tf.contrib.layers.xavier_initializer())  
b5 = tf.Variable(tf.random_normal([10]))  
hypothesis = tf.matmul(L4, W5) + b5
```

define cost/loss & optimizer

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=hypothesis, labels=Y))  
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

Epoch: 0001 cost = 0.266061549
Epoch: 0002 cost = 0.080796588
Epoch: 0003 cost = 0.049075800
Epoch: 0004 cost = 0.034772298
Epoch: 0005 cost = 0.024780529
Epoch: 0006 cost = 0.017072763
Epoch: 0007 cost = 0.014031383
Epoch: 0008 cost = 0.013763446
Epoch: 0009 cost = 0.009164047
Epoch: 0010 cost = 0.008291388
Epoch: 0011 cost = 0.007319742
Epoch: 0012 cost = 0.006434021
Epoch: 0013 cost = 0.005684378
Epoch: 0014 cost = 0.004781207
Epoch: 0015 cost = 0.004342310
Learning Finished!

Accuracy: **0.9742**

정답률 97.42%
→ overfitting이 발생한 듯!

dropout (keep_prob) rate 0.7 on training, but should be 1 for testing → Placeholder 0.7은 test 할때는 1로 바꿔 줘야 한다.

keep_prob = tf.placeholder(tf.float32)

W1 = tf.get_variable("W1", shape=[784, 512])

b1 = tf.Variable(tf.random_normal([512]))

L1 = tf.nn.relu(tf.matmul(X, W1) + b1)

L1 = tf.nn.dropout(L1, keep_prob=keep_prob)

W2 = tf.get_variable("W2", shape=[512, 512])

b2 = tf.Variable(tf.random_normal([512]))

L2 = tf.nn.relu(tf.matmul(L1, W2) + b2)

L2 = tf.nn.dropout(L2, keep_prob=keep_prob)

...

train my model

for epoch in range(training_epochs):

...

for i in range(total_batch):

batch_xs, batch_ys = mnist.train.next_batch(batch_size)

feed_dict = {X: batch_xs, Y: batch_ys, keep_prob: 0.7}

c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)

avg_cost += c / total_batch

Test model and check accuracy


correct_prediction = tf.equal(tf.argmax(hypothesis, 1), tf.argmax(Y, 1))

accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))

print('Accuracy:', sess.run(accuracy, feed_dict={

X: mnist.test.images, Y: mnist.test.labels, keep_prob: 1}))

Dropout for MNIST



Epoch: 0001 cost = 0.447322626
Epoch: 0002 cost = 0.157285590
Epoch: 0003 cost = 0.121884535
Epoch: 0004 cost = 0.098128681
Epoch: 0005 cost = 0.082901778
Epoch: 0006 cost = 0.075337573
Epoch: 0007 cost = 0.069752543
Epoch: 0008 cost = 0.060884363
Epoch: 0009 cost = 0.055276413
Epoch: 0010 cost = 0.054631256
Epoch: 0011 cost = 0.049675195
Epoch: 0012 cost = 0.049125314
Epoch: 0013 cost = 0.047231930
Epoch: 0014 cost = 0.041290121
Epoch: 0015 cost = 0.043621063
Learning Finished!

Accuracy: **0.9804!!**

Optimizers

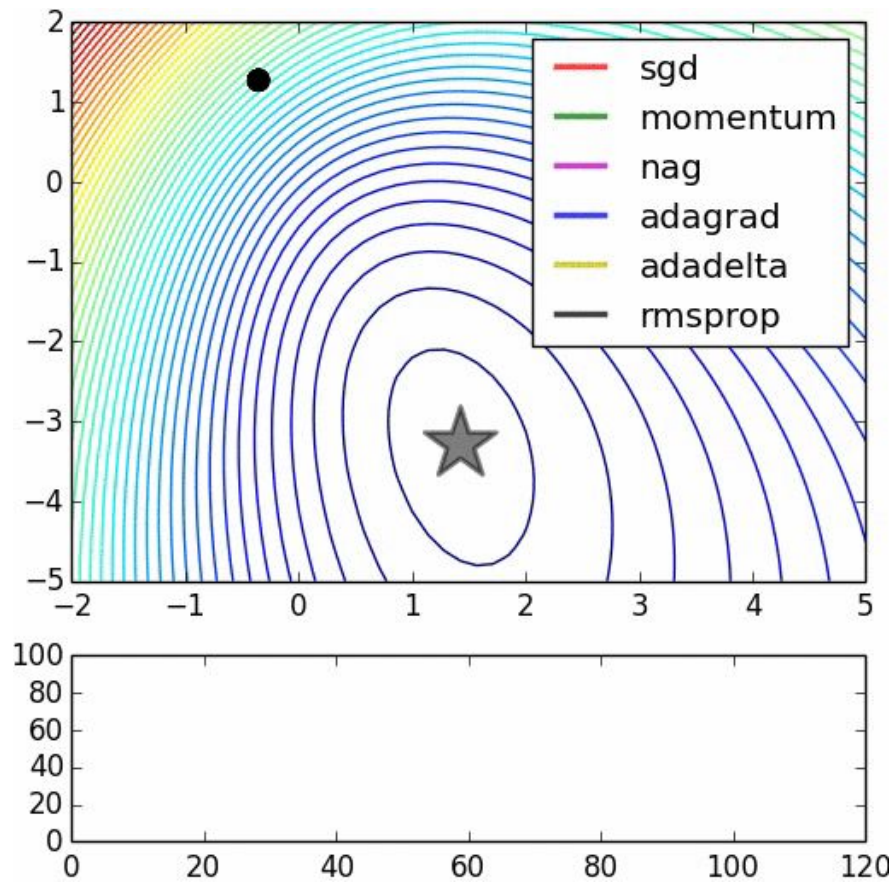
```
train = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```

Optimizers

```
train = tf.train.GradientDescentOptimizer(learning_rate=0.1).minimize(cost)
```


→ optimizer을 어떻게 줄이지 test+train을 같이 줄인다.

- `tf.train.AdadeltaOptimizer`
- `tf.train.AdagradOptimizer`
- `tf.train.AdagradDAOptimizer`
- `tf.train.MomentumOptimizer`
- `tf.train.AdamOptimizer`
- `tf.train.FtrlOptimizer`
- `tf.train.ProximalGradientDescentOptimizer`
- `tf.train.ProximalAdagradOptimizer`
- `tf.train.RMSPropOptimizer`



ADAM: a method for stochastic optimization

[Kingma et al. 2015]



일반적으로
Adam이 성능이 좋다
타 알고리즘에 비해

Use Adam Optimizer

```
# define cost/loss & optimizer
```

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(  
                                                                    logits=hypothesis, labels=Y))  
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

→ 사용법은 경사하강법과 동일하다.

Summary

- Softmax VS Neural Nets for MNIST, 90% and 94.5%

- Xavier initialization: 97.8%

- Deep Neural Nets ~~with~~ Dropout: **98%**

- Adam and other optimizers

- Exercise: Batch Normalization)

- https://github.com/hunkim/DeepLearningZeroToAll/blob/master/lab-10-6-mnist_nn_batchnorm.ipynb

→ 입력값을 normalize 잘하는 법.

Lecture and Lab I I

CNN

